

Sudoku validation as a pattern recognition problem

CHARU AGARWAL

Indian Institute of Technology, Dharwad
160010038@iitdh.ac.in

October 20, 2019

Abstract

This paper is an exploration of using neural networks to understand the constraints of knowledge games. We have considered the problem of checking whether a fully-filled sudoku board configuration is valid or invalid. The goal is to automatically infer the rules of sudoku from examples, without explicitly specifying them. We give an architecture of a model which may be suitable for this problem. While we have achieved good results for most errors, the model is not able to generalize for all cases. We consider some reasons when a problem may not be implicitly solvable by neural networks.

I. INTRODUCTION

Sudoku is a combinatorial [Lawler, 1985] number placement puzzle based on logical reasoning [Arnoldy]. The objective is to fill a 9x9 grid with digits so that each column, each row, and each of the nine 3x3 subgrids that compose the grid (also called "boxes", "blocks", or "regions") contain all of the digits from 1 to 9. The general sudoku problem is to reach a solution from a partially completed grid ¹. Figure 1 shows how a typical sudoku puzzle looks like.

Our goal is to train a neural network to understand the constraints of the sudoku. Particularly, given a board configuration, the model should be able to predict whether the configuration is valid or invalid. For the sake of simplicity, we have considered only fully-filled grids.

Now, one may argue the need of neural networks as this problem can be easily solved using an algorithm such as algorithm 1. Similar work has been done in [Grus, 2016], where the problem of "Fizz Buzz"² has been solved using neural networks with fair accuracy, but not reliably. The challenge is not to complicate simple problems by posing them as pattern recognition problems, but to develop a model which can automatically infer the rules of a game, without explicitly stating them. Simple problems can sometimes contain really interesting subtleties.

Sudoku, along with games such as crossword and wordsearch, roughly falls under the category of brain games which is an active area of research. By applying newly developing technologies in the areas of machine learning to understand how computers may be able to create such knowledge games, we hope to solve problems that have otherwise proven to be hard to solve using heuristic algorithms.

¹For a well-posed puzzle there exists a unique solution.

²The problem of fizzbuzz is defined as follows: print the numbers from 1 to 100, except that if the number is divisible by 3 print "fizz", if it's divisible by 5 print "buzz", and if it's divisible by 15 print "fizzbuzz".

Algorithm 1 Simple algorithm to classify a sudoku as valid or invalid.

```
1: procedure ISUDOKUVALID(grid)                                ▷ Validates a sudoku
2:   for each row do                                           ▷ Check for duplicate numbers in each row.
3:     for each number k in 1..n do
4:       if k is not in row then
5:         return False
6:       end if
7:     end for
8:   end for
9:   for each column do                                         ▷ Check for duplicate numbers in each column.
10:    for each number k in 1..n do
11:      if k is not in column then
12:        return False
13:      end if
14:    end for
15:  end for
16:  for each block do                                           ▷ Check for duplicate numbers in each block.
17:    for each number k in 1..n do
18:      if k is not in block then
19:        return False
20:      end if
21:    end for
22:  end for
23:  return True
24: end procedure
```

5	3			7				
6			1	9	5			
	9	8					6	
8				6				3
4			8		3			1
7				2				6
	6					2	8	
			4	1	9			5
				8			7	9

Figure 1: A typical sudoku puzzle.

II. THE DATASET

The sudoku puzzles for training have been taken from [Park, 2016]. The data contains 1 million sudoku games and solutions. The solutions are all valid solutions. Hence, for our purpose, we decide to generate invalid sudokus from the valid ones. In the generation of invalid dataset, for each invalid instance, the valid instance was included in the training data. This would make the learning difficult as close samples are in both positive and negative class.

The valid and invalid sudokus generated together form a dataset of 2 million samples. Depending of the method of generating invalid puzzles, we form three different datasets:

- Dataset A: In this dataset, 50% of the values (on an average) of the sudoku are in error. This has been achieved by shifting the value of some elements by a random integer between 1 - 8.
- Dataset B: In this dataset, one error has been introduced in the sudoku. This has been achieved by shifting the value of a random element by a random integer between 1 - 8.
- Dataset C: In this dataset, errors are introduced such that exactly one constraint (row/column/block) is not satisfied, while the other two constraints are satisfied. This is so that we get a more general distribution of errors. Figure 2 gives a schematic diagram of the transformations implemented.

The elements in the grid were standardized onto unit scale and labelled. The dataset after processing is shown in figure 3.

III. THE ARCHITECTURE

The performance of a model depends highly on the network topology, the learning law and the activation dynamics. Hence, we have carefully designed our model keeping in mind the advantages and disadvantages of particular neural networks. The architecture of our network is summarized in figure 4. We have three Convolutional Neural Network layers at the input, which have one filter each:

- Row sum layer: This layer consists of one 1x9 filter, with stride 1, to sum up the elements in each row.
- Column sum layer: This layer consists of one 9x1 filter, with stride 1, to sum up the elements in each column.
- Block sum layer: This layer consists of one 3x3 filter with stride 3 to sum up the elements in each block.

Convolutional neural network (CNN) is well suited for pattern classification and is designed specifically to recognize two-dimensional shapes with a high degree of invariance to translation, scaling, skewing, and other forms of distortion [Haykin, 2009]. The salient features of this model are local knowledge (each neuron takes inputs from a local receptive field in the previous layer, thereby forcing it to extract local features) and weight sharing (individual neurons are constrained to share the same weights). The rationale behind this choice is that if all the filter weights converge to 1, then the model should be able to detect anomalies in the filling (in most cases), as at least one of the filter outputs will not be equal to 4.5.³

Next, we have a layer which flattens and concatenates the outputs of the CNNs. The output vector is flattened and concatenated in the form $[x_1, x_2, \dots, x_{27}]$ where $x_1 - x_9$ are the outputs from

³Note that the sum conditions are necessary but not sufficient for a sudoku to be valid. For example, the sequence: 4 4 4 9 5 7 9 2 1 has sum 45.

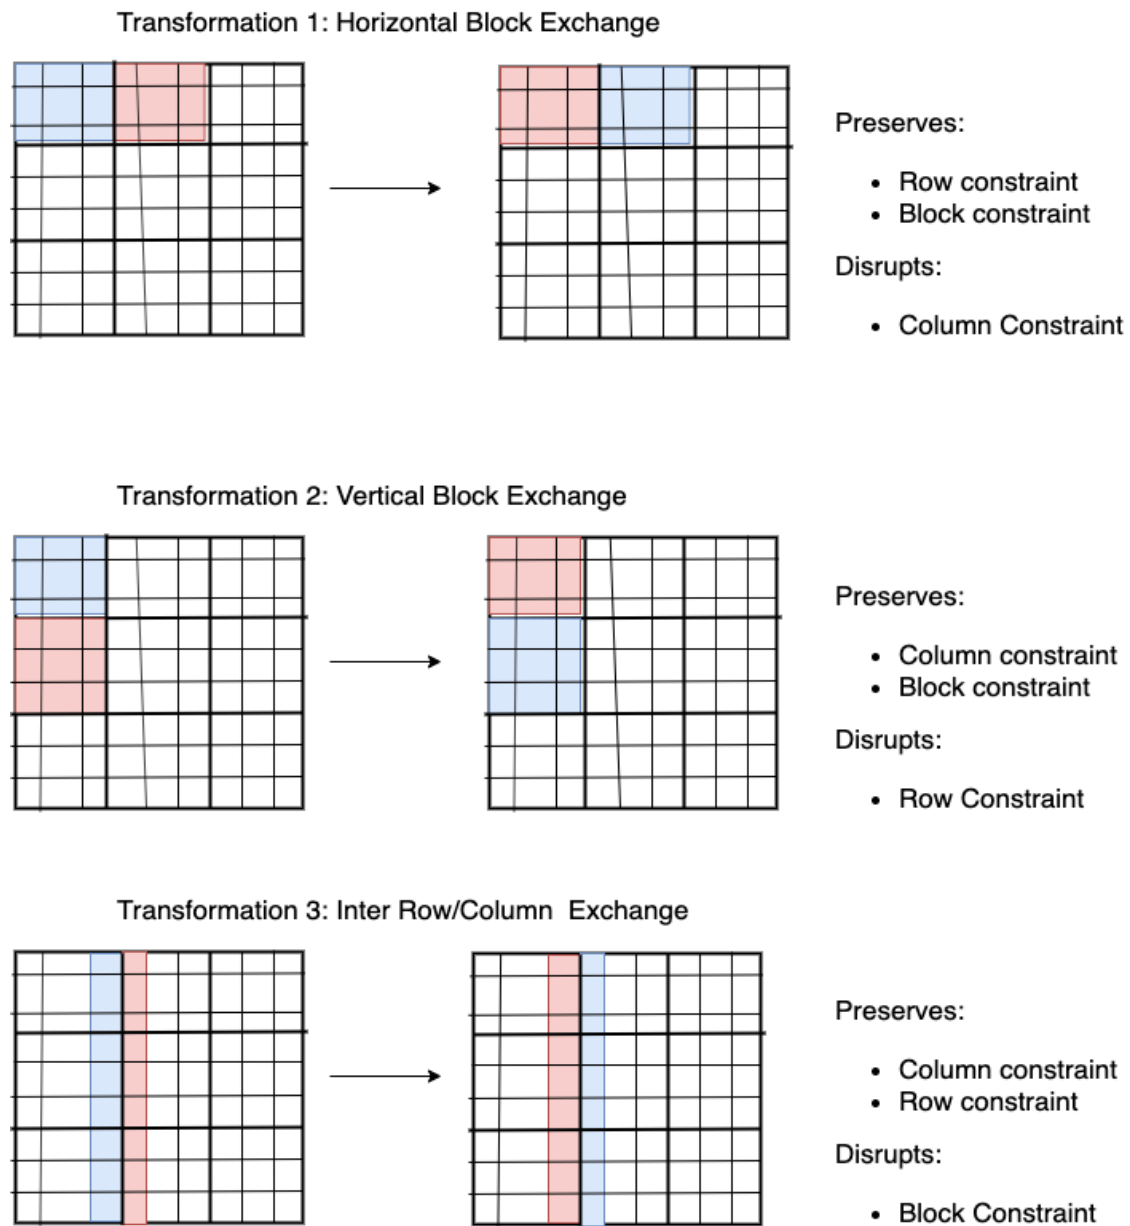


Figure 2: Transformations to convert a valid sudoku to invalid such that only one constraint is disrupted.

		Sudoku	Valid
0	[0.4, 0.3, 0.1, 0.9, 0.7, 0.2, 0.8, 0.5, 0.6, ...		0
1	[0.8, 0.4, 0.9, 0.3, 0.2, 0.6, 0.5, 0.1, 0.7, ...		0
2	[0.5, 0.6, 0.1, 0.2, 0.4, 0.7, 0.8, 0.9, 0.3, ...		1
3	[0.1, 0.4, 0.7, 0.9, 0.3, 0.2, 0.5, 0.6, 0.8, ...		0
4	[0.1, 0.7, 0.5, 0.3, 0.6, 0.8, 0.9, 0.4, 0.2, ...		1

Figure 3: Sample data used for training.

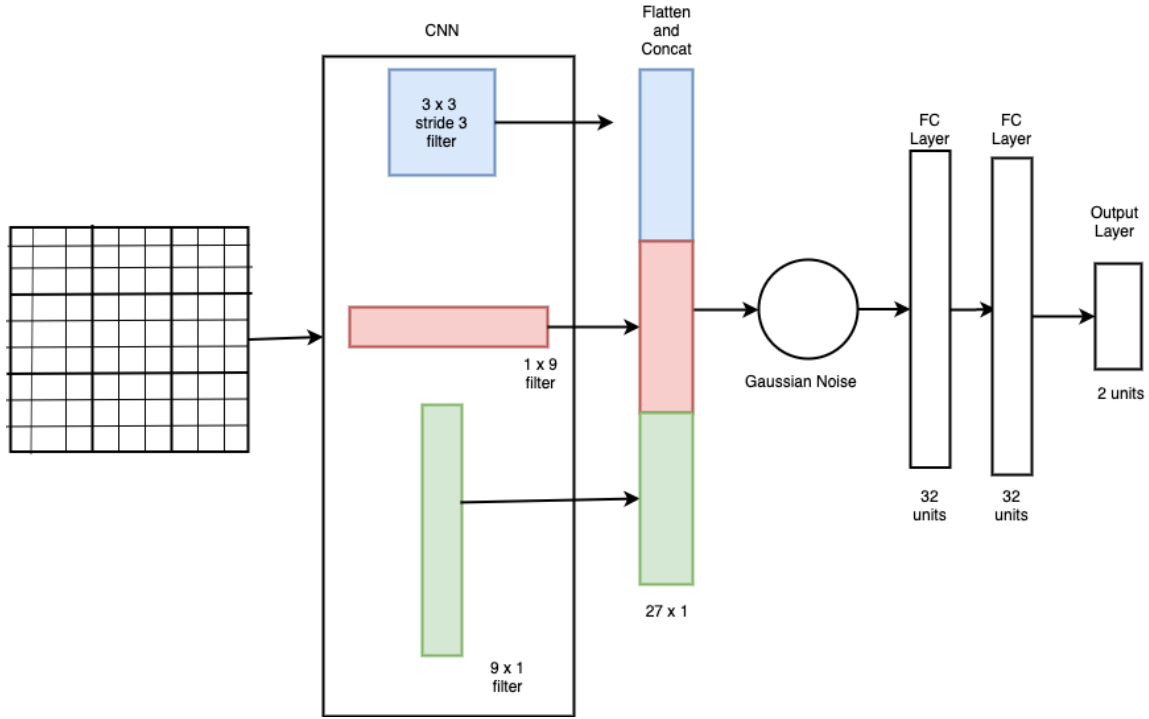


Figure 4: An illustration of the architecture of our network for classifying sudokus as valid and invalid.

block sum layer, $x_{10} - x_{18}$ are the outputs from row sum layer and $x_{19} - x_{27}$ are the outputs from the column sum layer. We add Gaussian Noise with mean 0 and standard deviation 0.01 to increase the spread of the valid class.⁴ Finally, the output was fed to a standard two-class classification network: MLP with relu transfer functions, 2 layers of 32 neurons each with Adam for optimization. A multilayer feedforward neural network with backpropagation learning on a finite set of independent and identically distributed samples leads to an asymptotic approximation of the underlying a posteriori class probabilities provided that the size of the training set data is large, and the learning algorithm does not get stuck in a local minimum [Hampshire and Pearlmutter, 1990].

IV. EVALUATION METRICS

We consider five standard metrics for evaluating the model:

- Accuracy: Accuracy is the fraction of predictions that are correct.

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN} \quad (1)$$

where TP = True Positives, TN = True Negatives, FP = False Positives, and FN = False Negatives.

- Loss: It is a summation of the errors made for each example in training or validation sets.

⁴The only example of the valid class is [4.5, 4.5, ..., 4.5]

- Precision: Precision is defined as the fraction of postive identifications that are actually correct.

$$Precision = \frac{TP}{TP + FP} \quad (2)$$

- Recall: Recall is the proportion of actual positives that are identified correctly.

$$Recall = \frac{TP}{TP + FN} \quad (3)$$

- AUROC: AUROC stands for "Area under the ROC Curve." An ROC curve (receiver operating characteristic curve) is a graph showing the performance of a classification model at all classification thresholds. AUROC can be interpreted as the probability that the model ranks a random positive example more highly than a random negative example.

V. RESULTS

First we fixed the weights of each filter to $[1, 1, \dots, 1]$ and allowed the MLP to learn the weights for classification. This gave us good accuracy ($> 96\%$). Hence, next we unclamped the network and allowed it to learn the filter weights too. Also, we seeded our experiments for reproducible results, and noticed the accuracy of the model varied strongly with the seed value.

Table 1 summarizes the results obtained from the model. While the model achieves excellent accuracy for dataset A and dataset B, it performs fairly well for dataset C. From the precision and recall, we can observe that both dataset A and C result in more false positives, indicating model is more likely to classify a sudoku as valid than invalid. The AUROC is 1 for datasets A and B, indicating that negative samples are ranked strictly below positive samples. For all the datasets, the AUROC is more than 0.5, which means the model performs significantly better than a random classifier.

Table 1: Results for sudoku classification

Dataset	Test Accuracy	Test Loss	Precision	Recall	AUROC
A	0.99	1.86e-05	0.99	1.00	1.00
B	1	6.56e-07	1.00	1.00	1.00
C	0.72	0.50	0.64	0.99	0.73

VI. DISCUSSION

In this section, we try to understand why the model works for most cases, and more importantly, why it is not able to classify a general problem. To understand why the model is able to classify well for dataset A and dataset B, we decided to look at the filter weights learnt by the model. Figure 5 shows the filter weights before and after learning. The model has learnt the column sum filter as $c * [1, 1, \dots, 1]$ (for some constant c), which is able to detect most errors. We observe that not much learning has occurred for the other two filters. This will not work for the general case where the errors may be such that the net sum of errors along a column is zero, although the errors themselves are non zero.

On intializing the filter weights with Gaussian Noise with mean 1 and standard deviation 0.01, all the filter weights converge to $c * [1, 1, \dots, 1]$ (for some constant c). However, if we increase the




Figure 5: Weights before (left) and after (right) learning. Most of the learning appears to happen in the column sum filter where the weights are all 1.4^* . Hence, we can say this filter has understood the constraints as it is assigning almost equal weights to each element. This can explain the good performance of the model. (For most errors, the column filter should be able to detect errors. This will not work well for multiple errors where the net sum is 0 within same column.)




Figure 6: Visualization of Dataset A (puzzles when 50% errors are introduced (on an average) in invalid sudokus): PCA (left) vs t-SNE (right)

variance of the Gaussian Noise to 1, the filter weights are not learnt. We had considered the possibility that only one filter weights are being learnt because it is not necessary for the model to learn the other two filters, since one filter is sufficient to detect most errors.⁵ However, this hypothesis was disproved as we tried to train the model with only the block sum filter, but still it was unable to learn the weights.

We decided to visualize the datasets using dimensionality reduction techniques: Principal Component Analysis (PCA) and t-Distributed Stochastic Neighbouring Entities (t-SNE). Figures 6, 7 and 8 show the plots for the datasets. We found that while the valid and invalid classes of dataset A are well-separated, classes of dataset B and C are not easily separable. This may suggest why the model is not able to generalize, since the data itself is inherently not separable into two classes. Overall, t-SNE is able to better cluster the examples into the two classes as compared to PCA.

⁵Part of the reason why dataset C was introduced, since it disrupts only one constraint at a time.




Figure 7: Visualization of Dataset B (puzzles when 1 error is introduced in invalid sudokus): PCA (left) vs t-SNE (right)




Figure 8: Visualization of Dataset C (puzzles with one constraint disrupted in invalid sudokus): PCA (left) vs t-SNE (right)

VII. CONCLUSIONS

We considered the idea of modelling the task of validating a completely filled sudoku board configuration as a pattern recognition problem. That is, we wanted to see that given a lot of examples from the valid and invalid classes, can we train a neural network to learn the constraints of sudoku. We proposed a CNN-MLP based model architecture and we observed that while the model is able to classify most samples as valid or invalid correctly, it has only partially learnt the constraints. The good accuracy is attributed to the learning of the column sum filter weights, which allows the model to sum up the elements in each column and detect most forms of discrepancies.

We conclude that while we have partially solved the problem of sudoku classification, solving the full problem may require a different model architecture altogether. Observing the data, it seems difficult to model this problem as a pattern recognition problem since the distribution of the valid and invalid class are not well-separated. Our model was quite tailored for the 9x9 sudoku puzzles, but shall not be able to generalize for sudokus of different sizes. But all is not lost, since this might be a small step for sudoku classification, but a significant step for understanding knowledge games. Modeling the problem of learning of the constraints of knowledge games as a pattern recognition problem is a powerful idea, since now we do not need to tell the machine the rules of the games.

VIII. ACKNOWLEDGEMENTS

This study is the result of many fruitful discussions with Prof. Sudheendra Hangal and Prof. Prabuchandran KJ, who have been the advisors of the project. The author would also like to thank Dept. of Computer Science, IIT Dharwad, for providing the opportunity to pursue this problem. This project has been a great learning experience on machine learning problems and has introduced the author to many techniques to design, analyze and debug neural networks.

REFERENCES

- [Park, 2016] Kyubyong Park (2016). 1 million Sudoku games, Version 3. Retrieved October 6, 2019 from <https://www.kaggle.com/bryanpark/sudoku>.
- [Grus, 2016] Joel Grus (2016). Fizz Buzz in Tensor Flow. <https://joelgrus.com/2016/05/23/fizz-buzz-in-tensorflow/>.
- [Arnoldy] Ben Arnoldy. Sudoku Strategies *The Christian Science Monitor*.
- [Lawler, 1985] E. L. Lawler (1985). The Traveling Salesman Problem: A Guided Tour of Combinatorial Optimization. West Sussex: John Wiley & Sons. ISBN 0-471-90413-9
- [Haykin, 2009] S. S. Haykin (2009). Neural networks and learning machines/Simon Haykin. New York: Prentice Hall
- [Hampshire and Pearlmutter, 1990] J. B. Hampshire, and B. Pearlmutter (1990). Equivalence proofs for multilayer perceptron classifiers and the Bayesian discriminant function. *E InD. S. Touretzky, JL Elman, TJ Sejnowski, and GE Hinton, editors. Proceedings of the 1990 Connectionist Models Summer School. Morgan Kaufmann, San Mateo, CA.*